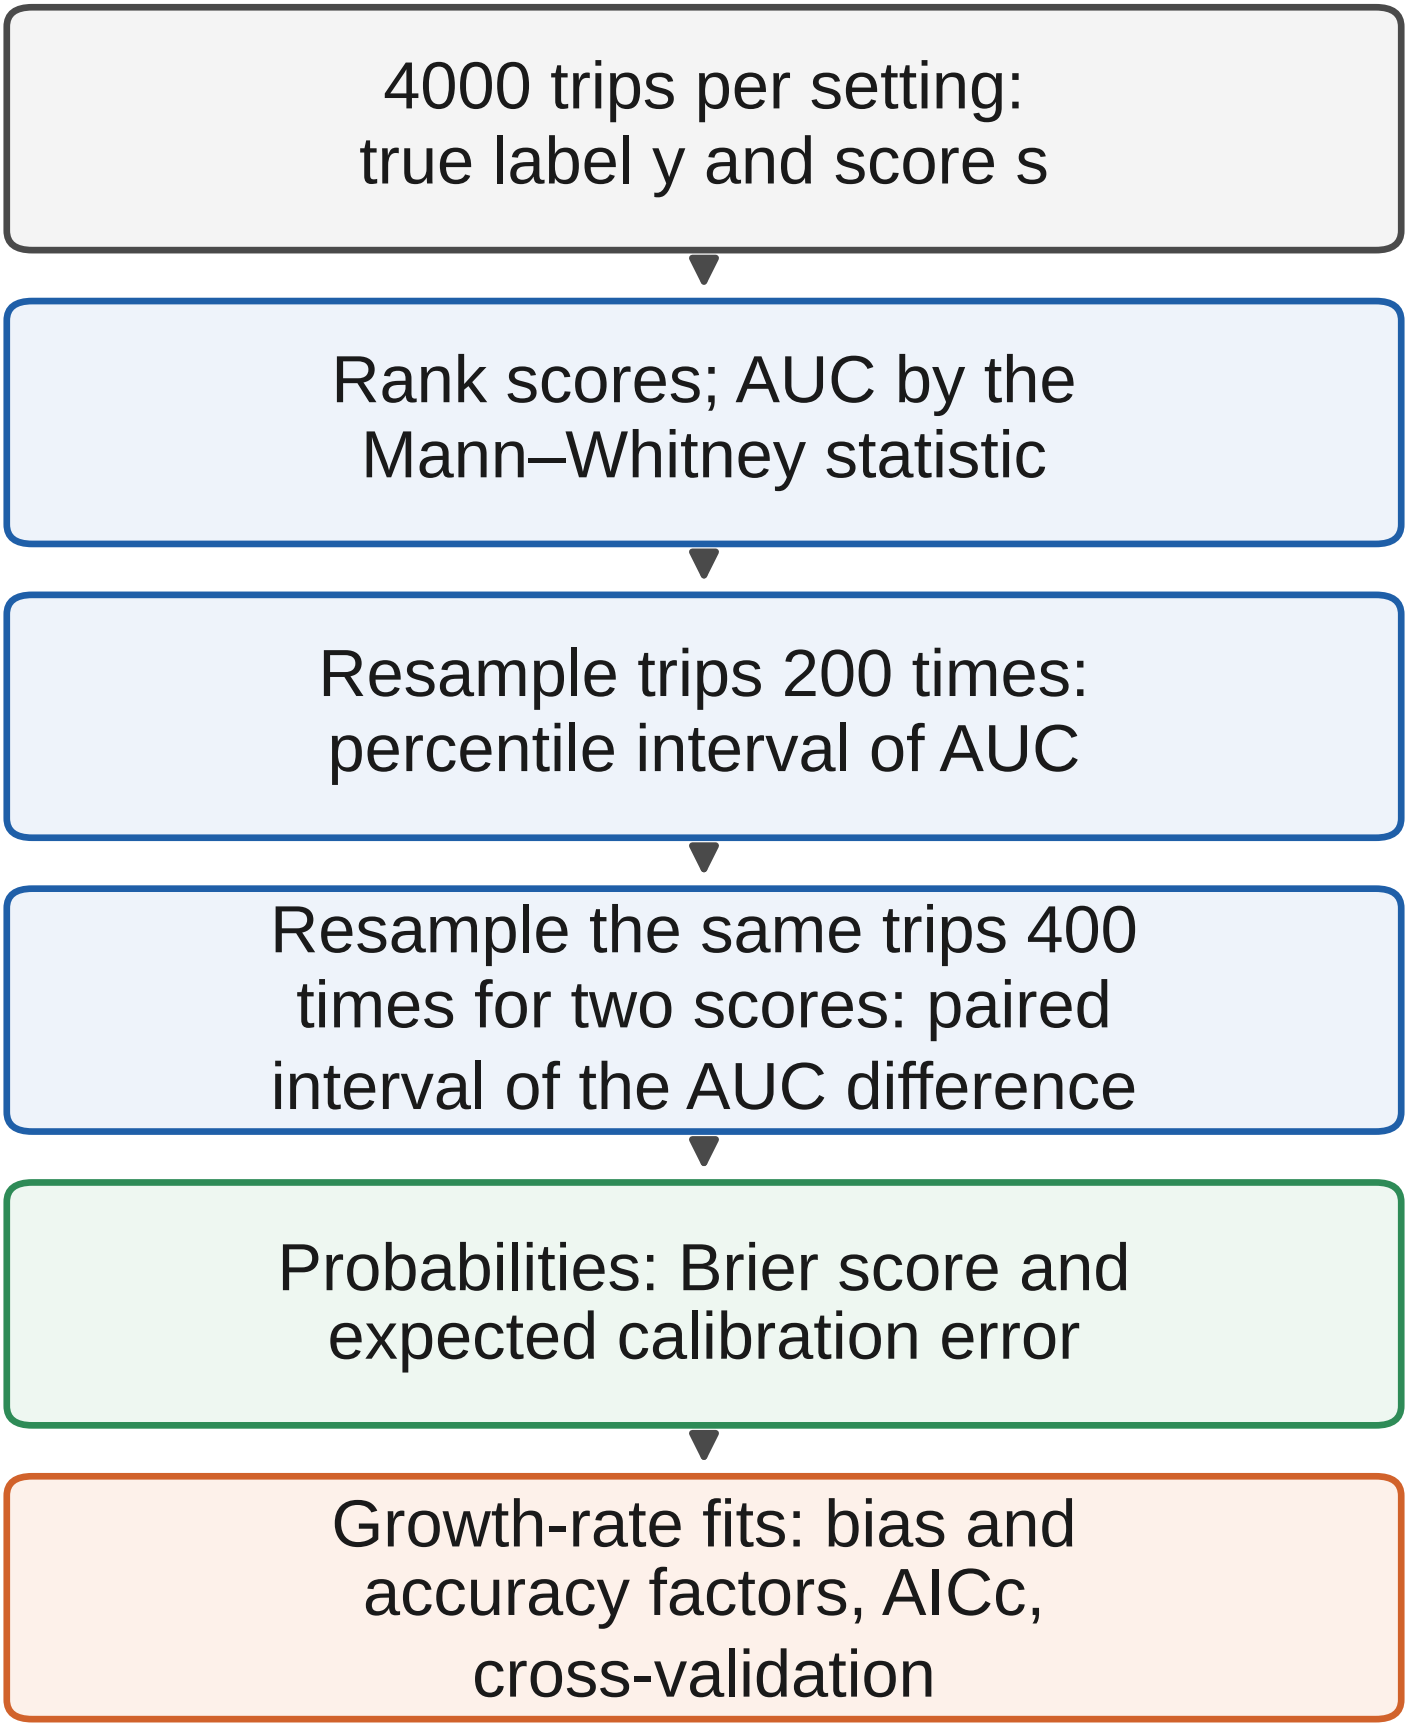
4000 trips per setting:
true label y and score s
Rank scores; AUC by the
Mann–Whitney statistic
Resample trips 200 times:
percentile interval of AUC
Resample the same trips 400
times for two scores: paired
interval of the AUC difference
Probabilities: Brier score and
expected calibration error
Growth-rate fits: bias and
accuracy factors, AICc,
cross-validation